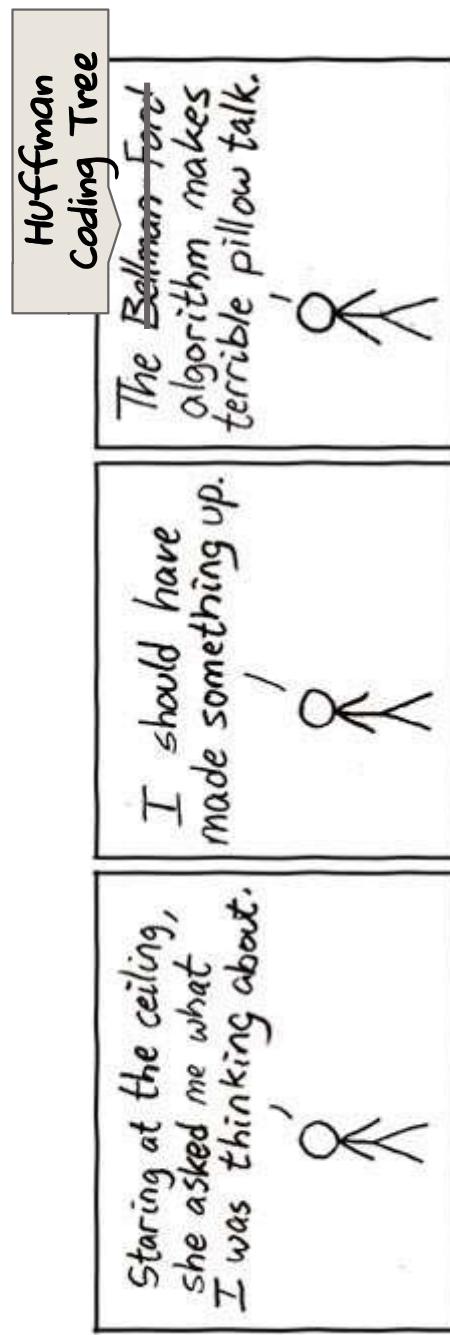


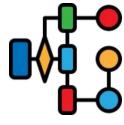
# Optimization Methods Part 1

## Greedy algorithms



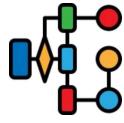
<https://xkcd.com/69/>

Richard Lobb and R Mukundan  
Department of Computer Science and Software Engineering  
University of Canterbury



## The Greedy Method ("Greedy heuristic")

- A simple *intuitive* method.
- Choose whatever option looks best at the moment
  - Hoping that a locally optimal choice will lead to a globally optimal solution.
- A decision made in one stage is not changed later.
  - No backtracking
- Greedy algorithms do not always yield optimal solutions, but for many problems, they do.
- In each iteration, greedy algorithms seek a *feasible* solution that minimizes or maximizes a given *objective* function.



## Examples you know and love already

### Minimum Spanning Tree algorithms:

- Prim's algorithm

- Maintain set of connected nodes
- Repeatedly add "cheapest" node

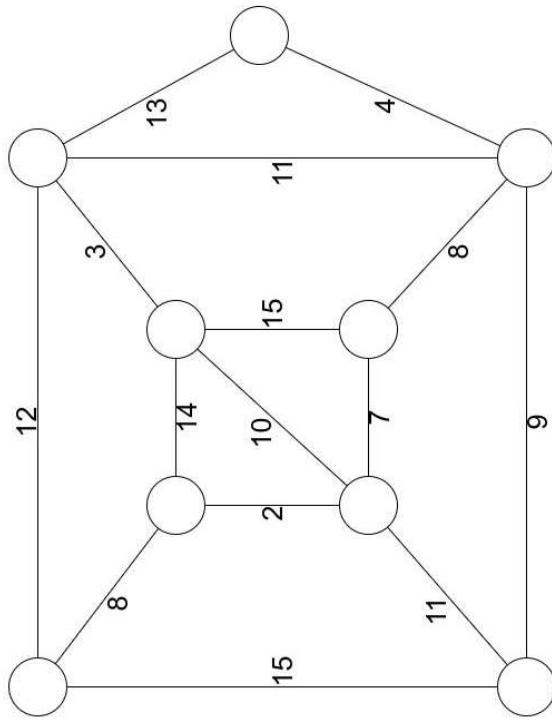
- GREEDY!

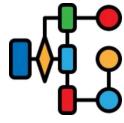
- Kruskal's algorithm

- Maintain forest of disconnected trees
- Repeatedly add "cheapest" edge that doesn't cause a cycle

- GREEDY!

Both fairly easily shown to be **optimal**

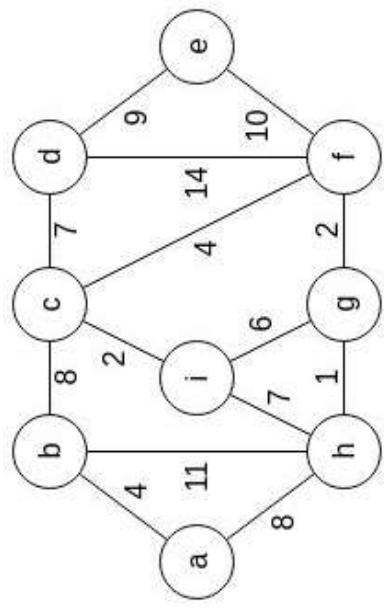




## Kruskal's Algorithm

**edges sorted by weight:**

1:gh, 2:ci, 2:fg, 4:ab, 4:cf, 6:gi, 7:cd, 7:hi,  
8:ah, 8:bc, 9:de, 10:ef, 11:bh, 14:df



**forest = set()**

**for each  $v \in V$ :**

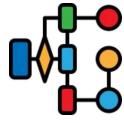
create new component {v}

**for each  $(u,v) \in E$  in order of non-decreasing weight:**

**if  $u$  and  $v$  are in different components:**

**forest.add((u,v))**

**merge components of  $u$  and  $v$**

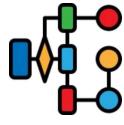


# The Coin Changing Problem



Image taken from Bioinformatics Algorithms, MIT Press, 2004.

You are a shopkeeper. You need to give change of  $V$  cents to a customer. The till has coins of denominations  $C = \{c_1, \dots, c_m\}$ . The goal is to give the customer the minimum number of coins whose total value equals the amount  $V$ .



# The Coin Changing Problem

## Assumptions:

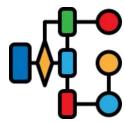
- $c_1 < c_2 < \dots < c_m$  (Input sorted in ascending order)
- Each denomination is available in unlimited quantity
- $c_1 = 1$ , so that a solution exists for any integer  $V > 0$

## Required solution:

Integers  $n_1, n_2, \dots, n_m$  such that

$$V = n_1 c_1 + n_2 c_2 + \dots + n_m c_m$$

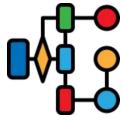
$N = n_1 + n_2 + \dots + n_m$  is the minimum possible value.



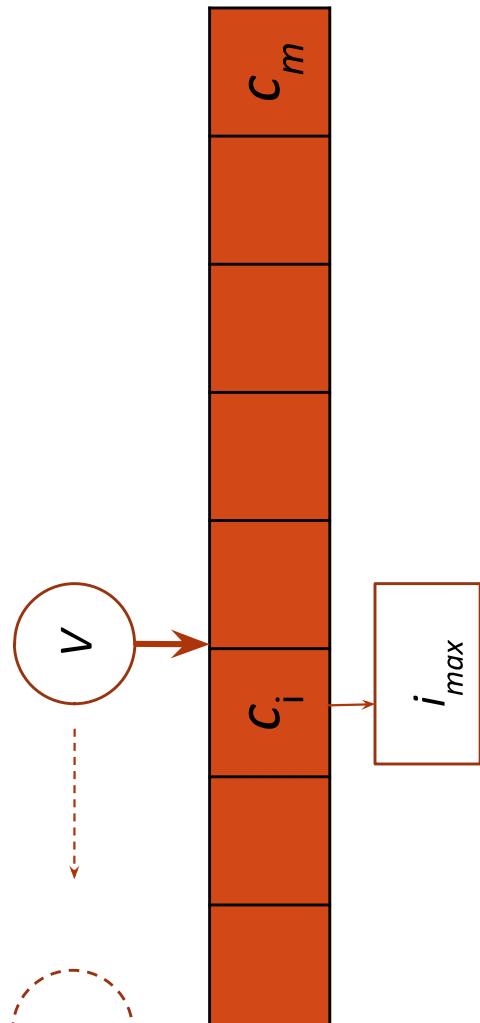
# The Coin Changing Problem

## Greedy strategy:

```
counts = defaultdict(int) # Counters of all coins
while V > 0:
    select the largest value  $c_i \leq V$  (i.e.,  $c_i \leq V < c_{i+1}$ )
    counts[ci] += 1
    V = V - ci
```



# The Coin Changing Problem



```
Sort C in ascending order
counts = defaultdict(int) # Coin counters
```

$i_{\max} = m$

**while**  $V > 0$ :

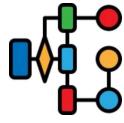
**while**  $c_{i_{\max}} > V$ :

$i_{\max} = i_{\max} - 1$

$\text{counts}[c_{i_{\max}}] += 1$

$V = V - c_{i_{\max}}$

Q: what precondition(s) must  
be satisfied to ensure these  
loops terminate?



# The Coin Changing Problem

Examples:

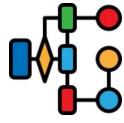
- $V = 82c \quad C = \{1c, 5c, 10c, 25c\}$

Greedy Solution:  $3*25c + 1*5c + 2*1c \quad (N = 6 \text{ coins})$

Here, the solution is optimal.

- $V = 80c \quad C = \{1c, 10c, 25c\}$
- Greedy Solution:  $3*25c + 5*1c \quad (N = 8 \text{ coins})$
- This is not an optimal solution.
- Optimal Solution:  $2*25c + 3*10c \quad (N = 5 \text{ coins})$

Can you suggest a sufficient property of the currency system that ensures optimality of a greedy algorithm for all  $V \geq 0$ ?



## Activity-Selection Problem

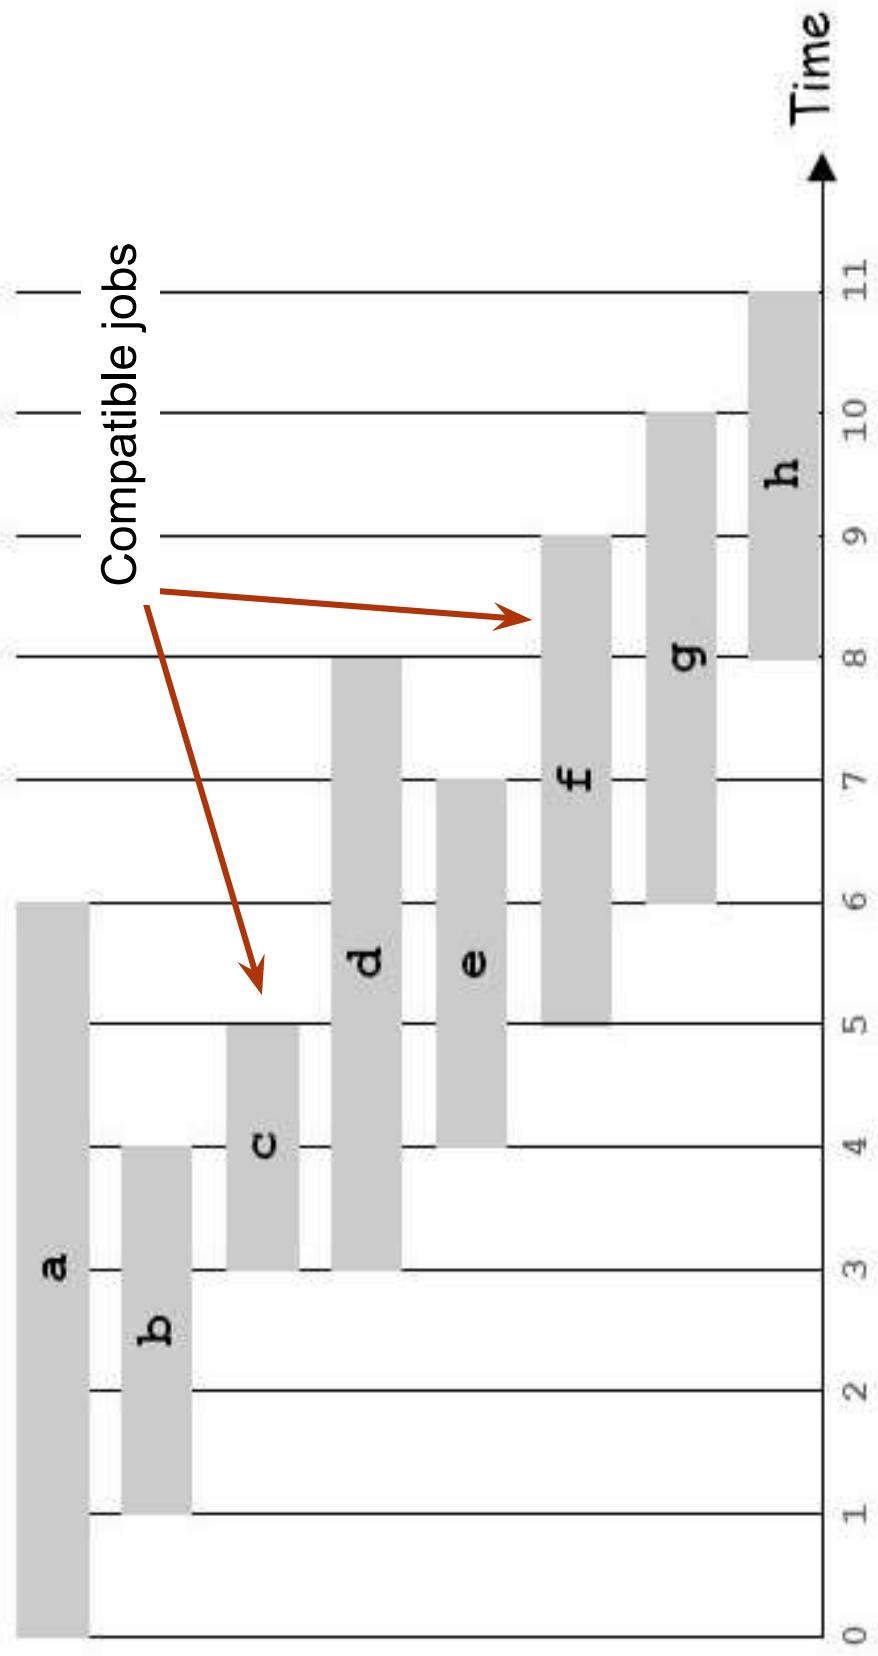
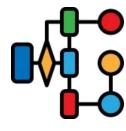
How to get your money's worth out of a carnival?

- Buy a wristband that lets you onto any ride
- Lots of rides, each starting and ending at different times
- Your goal: ride as many rides as possible

## The Specific Problem: Interval Scheduling

- Given a set of scheduled jobs
- Job  $j$  starts at time  $s_j$  and finishes at time  $f_j$
- Two jobs are compatible if they don't overlap
- Your goal: find maximum subset of mutually compatible jobs

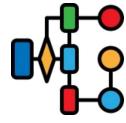
# Interval Scheduling



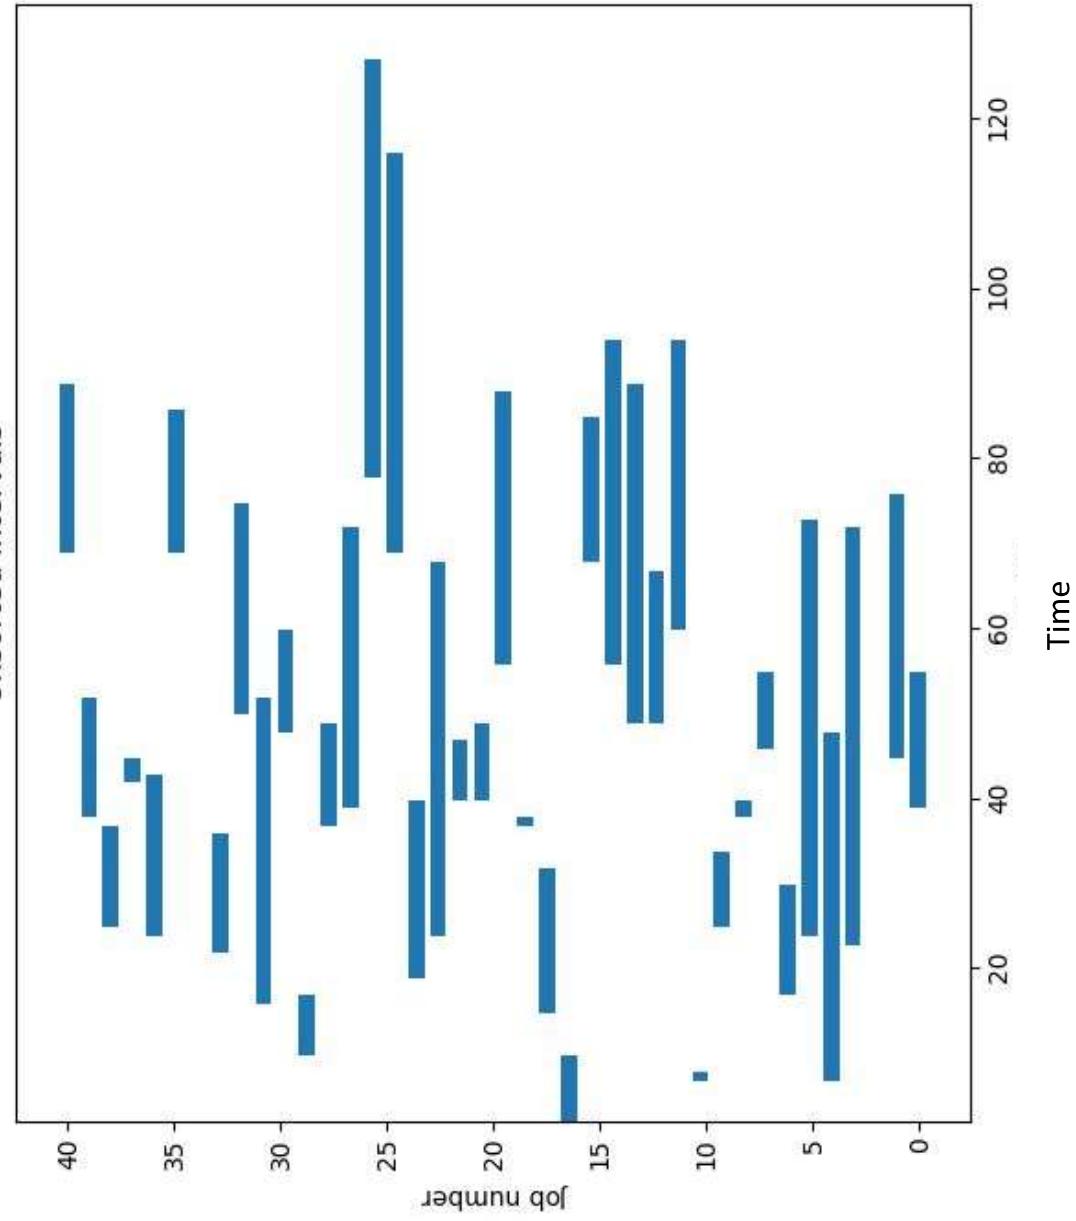
What is the (or an) optimal solution? Answer: {b, e, h}

11

# A more challenging example



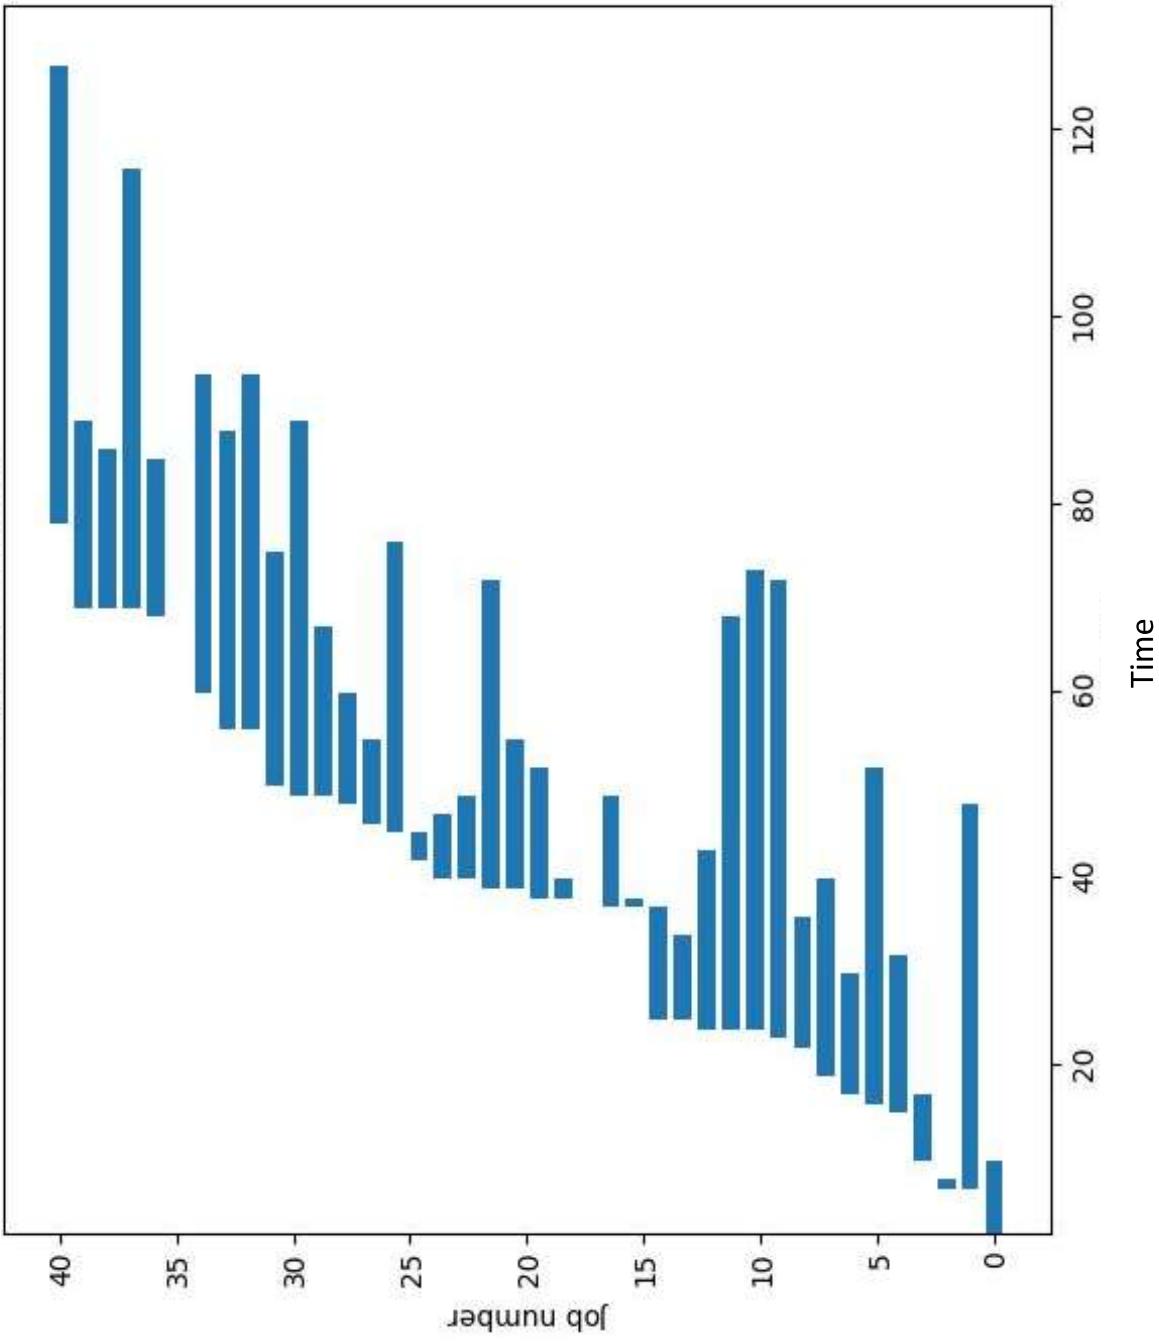
Unsorted intervals



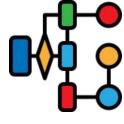
What algorithm  
will you use?

Assume greedy  
heuristic ...  
what is the sort  
key?

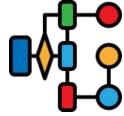
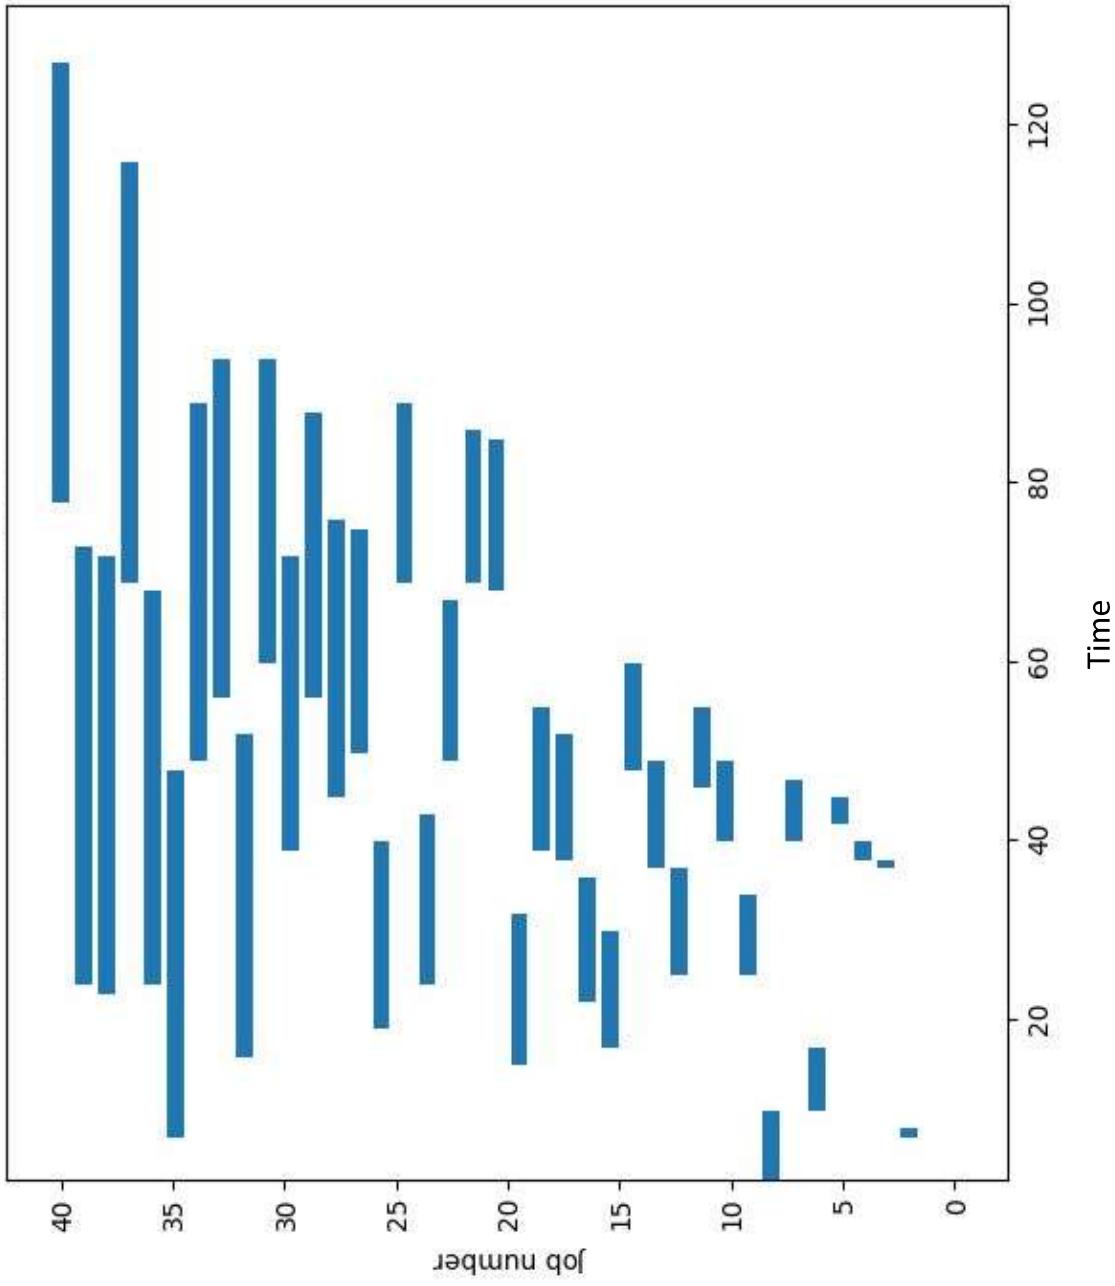
# Sort by start time?

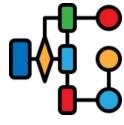


13

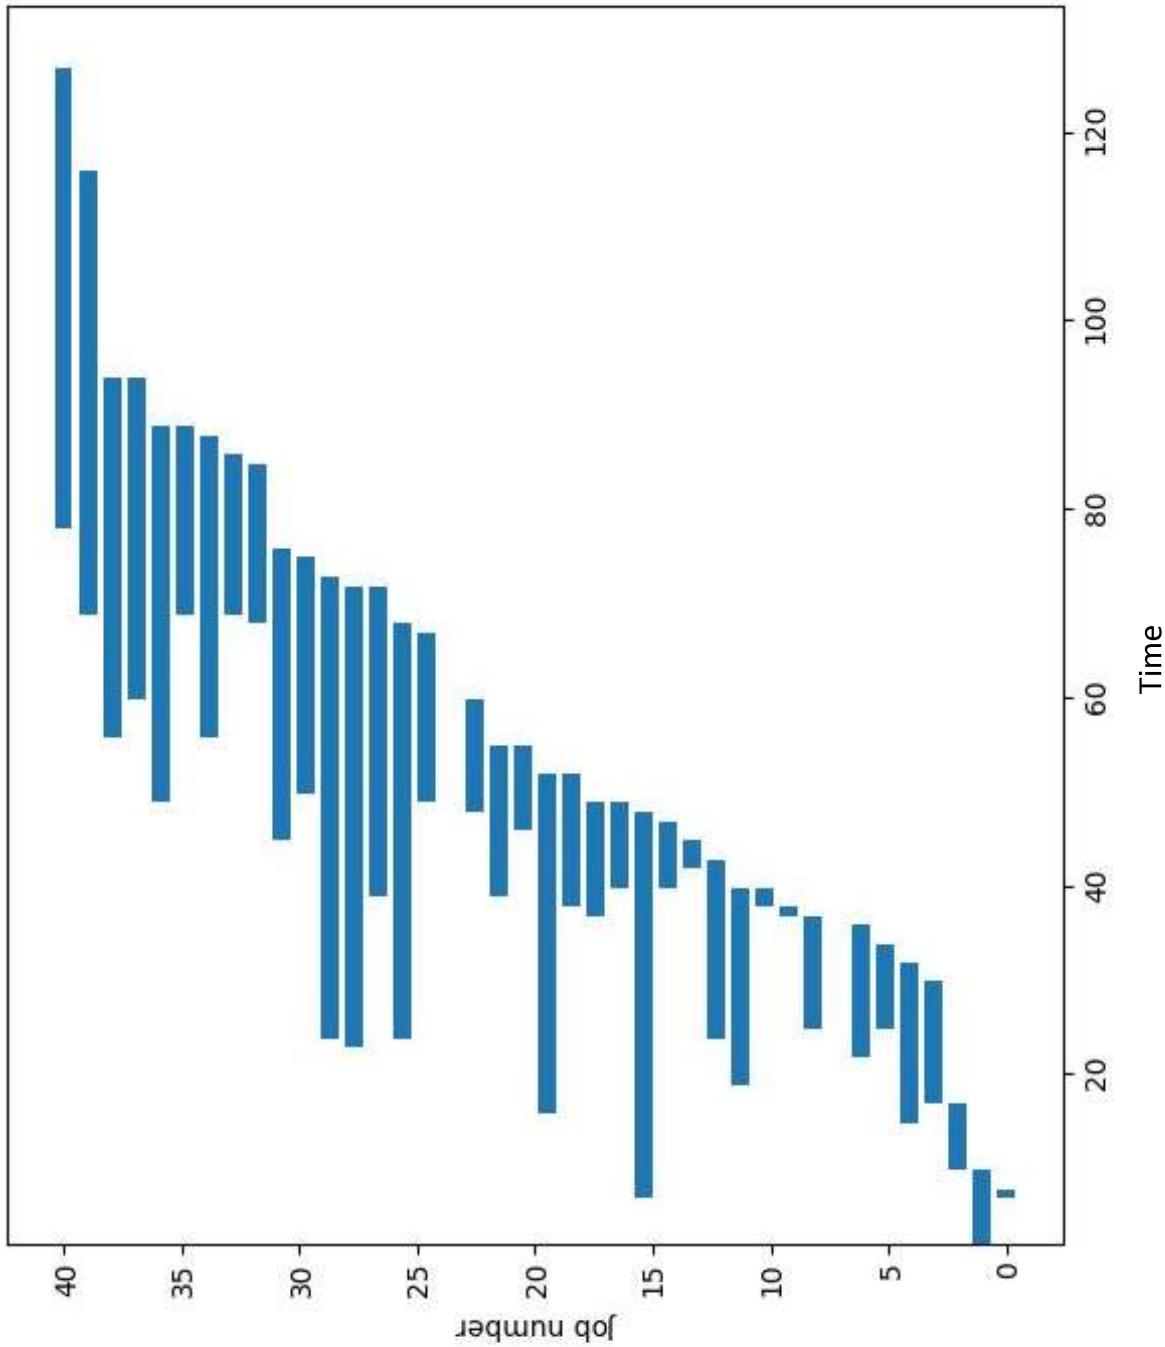


# Sort by duration?

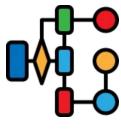




Sort by end time?

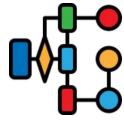


15



## Interval Scheduling

- Have to *order* the jobs. Then take each job, provided it is compatible with the ones already taken.
- Bad choices for ordering:
  - Increasing order of start times  $s_j$ 
    - "Oh no, the first job takes all day"
  - Increasing order of job intervals  $f_j - s_j$ 
    - "Oh no, the shortest job overlaps two mutually compatible jobs"
- The right choice for ordering:
  - Increasing order of finish times  $f_j$
  - On completion of job with finish time  $f_j$ , start next job ( $s_k, f_k$ ) in the sorted order such that  $s_k \geq f_j$ .



# Interval Scheduling Algorithm

## Greedy algorithm:

Sort jobs by finish times:  $f_1 \leq f_2 \leq \dots \leq f_n$

$S = \{\}$  (Set of jobs selected)

$t_{\text{current}} = 0$  (Finish time of the last job selected)

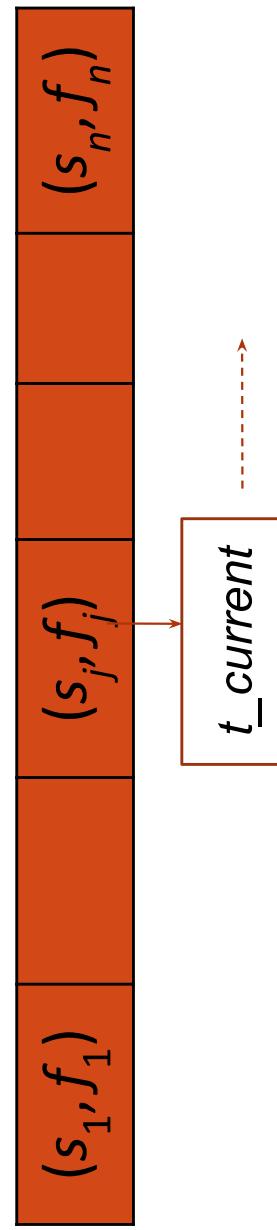
**for**  $j = 1..n$ :

**if**  $s_j \geq t_{\text{current}}$  :

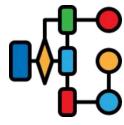
$S = S \cup \{j\}$

$t_{\text{current}} = f_j$

**return**  $S$



Is this always optimal? If so, prove it. If not, give a counter-example.



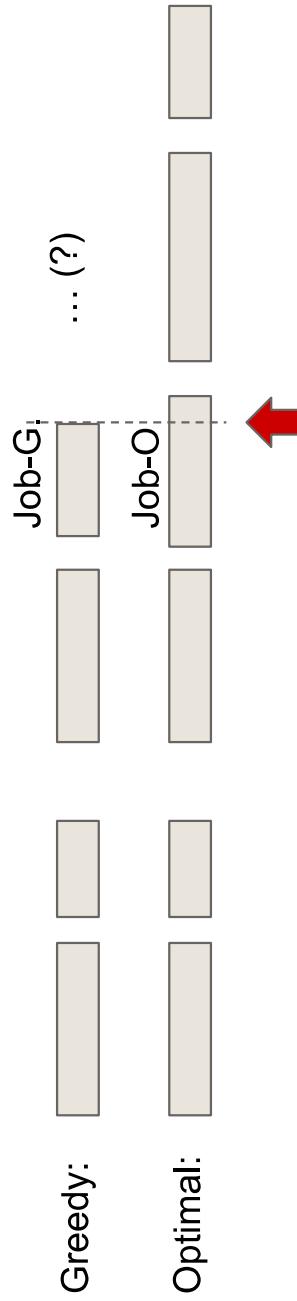
# Interval Scheduling Optimality

Yes, greedy algorithm (sorting by finish time) *is* optimal.

**Proof by contradiction (informal):**

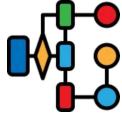
Assume algorithm is **not** optimal. Then for some problem there exists a better solution than that found by the greedy algorithm.

Consider the point at which the results first differ non-trivially (e.g. excluding simple swapping of one job for another).



Optimal algorithm has Job-O that finishes later than Job-G. But if optimal algorithm had instead used Job-G the rest of its jobs could remain unchanged and it would still be optimal. Therefore this **isn't** the point at which the results first differ non-trivially, contradicting the assumption.

More formally: see [https://en.wikipedia.org/wiki/Charging\\_argument](https://en.wikipedia.org/wiki/Charging_argument)

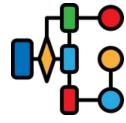


## Fractional Knapsack Problem

- Given a set  $S$  of  $n$  items, such that each item  $i$  has a positive value  $b_i$ , and a positive weight  $w_i$ , determine the subset of  $S$  that yields the maximum value without exceeding a given weight  $W$ .
- We can select a fraction  $x_i$  of an item  $i$  giving a value  $x_i b_i$  and a weight  $x_i w_i$ .

$$W = 9$$

<b><math>i</math></b>	<b>Item</b>	<b>Benefit <math>b_i</math></b>	<b>Weight <math>w_i</math></b>
1	Popcorn (1 carton)	12	4
2	Potato chips (1 bag)	15	3
3	Pizza (1)	14	2
4	Chocolate cookies (bag)	20	5



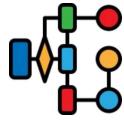
# Fractional Knapsack Problem

Greedy Solution with measure  $b_i$

<b><math>i</math></b>	<b>Item</b>	<b><math>x_i</math></b>	<b>Benefit <math>x_i b_i</math></b>	<b>Weight <math>x_i w_i</math></b>
4	Chocolate cookies	1	20	5
2	Potato chips	1	15	3
3	Pizza	$\frac{1}{2}$	7	1
	TOTAL		42	9

Greedy Solution with measure  $b_i / w_i$

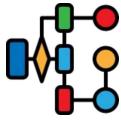
<b><math>i</math></b>	<b>Item</b>	<b><math>x_i</math></b>	<b>Benefit <math>x_i b_i</math></b>	<b>Weight <math>x_i w_i</math></b>
3	Pizza	1	14	2
2	Potato chips	1	15	3
4	Chocolate cookies	$\frac{4}{5}$	16	4
	TOTAL		45	9



# Fractional Knapsack Problem

## Correct greedy algorithm (optimal)

- Sort items in descending order of  $b_i/w_i$
- Take items (or a fraction of the final item) in order until knapsack full.
- Gives an optimal solution in  $O(n \log n)$  time.



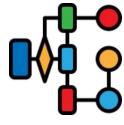
## The 0/1 Knapsack

(not amenable to a greedy solution)

- A thief robbing a store finds  $n$  items; the  $i^{th}$  item is worth  $b_i$ , dollars, and weighs  $w_i$  kgs ( $b_i, w_i$  are integers). The thief can carry at most  $W$  kgs. Which items should (s)he take for maximum payload?
- A greedy approach is in general sub-optimal.
- Example:

	(a)	(b)	(c)	(d)	(e)	
	3 kg	3 kg	4 kg	4 kg	8 kg	10 Kgs
	\$45	\$45	\$80	\$80	\$100	

- Greedy solution using only  $b_i$ : (e) = \$100  
or using  $b_i/w_i$ : (c) + (d) = \$160
- Optimal solution: (a) + (b) + (c) = \$170  
• Needs dynamic programming - next week.



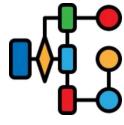
## Representing text in binary

The ASCII encoding system uses a fixed-length 7-bit binary string to store each character (but using 8-bit bytes)

- ASCII = American Standard Code for Information Interchange

0100000	0110000	0	1000000	@	1010000	P	1100000	`	1110000	p
0100001	0110001	1	1000001	A	1010001	Q	1100001	a	1110001	q
0100010	0110010	2	1000010	B	1010010	R	1100010	b	1110010	r
0100011	0110011	3	1000011	C	1010011	S	1100011	c	1110011	s
0100100	0110100	4	1000100	D	1010100	T	1100100	d	1110100	t
0100101	0110101	5	1000101	E	1010101	U	1100101	e	1110101	u
0100110	0110110	6	1000110	F	1010110	V	1100110	f	1110110	v
0100111	0110111	7	1000111	G	1010111	W	1100111	g	1110111	w
0101000	0111000	8	1001000	H	1011000	X	1101000	h	1111000	x
0101001	0111001	9	1001001	I	1011001	Y	1101001	i	1111001	y
0101010	0111010	:	1001010	J	1011010	Z	1101010	j	1111010	z
0101011	0111011	,	1001011	K	1011011	[	1101011	k	1111011	{
0101100	0111100	<	1001100	L	1011100	\`	1101100	l	1111100	\`
0101101	0111101	=	1001101	M	1011101	]	1101101	m	1111101	}
0101110	0111110	>	1001110	N	1011110	^	1101110	n	1111110	~
0101111	0111111	?	1001111	O	1011111	1101111	o	1101111	o	1111111

For a more complete table, see <https://www.ascii-code.com/>



# Fixed-length versus variable length codes

## International Morse Code

Fixed-length codes (e.g. 8-bit ASCII) are convenient:

- Easy indexing, counting etc
- But variable-length codes are (potentially) more efficient
- Example 1: Morse Code

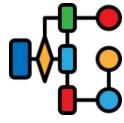
International Morse Code

1. The length of a dot is one unit.
2. A dash is three units.
3. The space between parts of the same letter is one unit.
4. The space between letters is three units.
5. The space between words is seven units.

U	• • -
V	• • • -
W	• - -
X	- - .
Y	- - :.
Z	- - -

A	• -
B	• • -
C	• • • -
D	• - -
E	.
F	• • - -
G	• - - -
H	• - - -
I	• •
J	• - - -
K	• - - -
L	• - - -
M	• - -
N	• - -
O	• - - -
P	• - - -
Q	• - - -
R	• - - -
S	• • •
T	-

- Most-frequent letters have short encodings
- Example 2: UTF-8
- Most-common chars are 1-byte, less common 2, 3 and 4 bytes



## An example problem

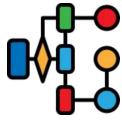
Consider the following string of characters containing only the six symbols ‘a’, ‘b’, ‘c’, ‘d’, ‘e’, ‘f’:

**acbfcffffccadbbcccdceadeaaaaacbccbebefbea**

The character frequencies are:

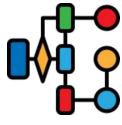
Char	Frequency	No. of bits (8-bit ASCII)
‘a’ = 01100001	9	72
‘b’ = 01100010	8	64
‘c’ = 01100011	15	120
‘d’ = 01100100	3	24
‘e’ = 01100101	5	40
‘f’ = 01100110	2	16
<b>TOTAL</b>	<b>42 characters</b>	<b>42 * 8 = 336 bits</b>

How to encode this “document” efficiently?



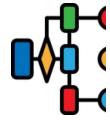
## Variable length binary encoding

- How might we improve that by using variable length encodings?
- Ambiguous decoding:
  - ‘a’=0, ‘b’=1, ‘c’=01, ‘d’=10, ‘e’=11, ‘f’=001
  - 010 could mean either “aba”, or “ad”, or “ca”
  - We just have a stream of 0s and 1s, no spacers between them (c.f. Morse).
- A binary code allows for unambiguous decoding if and only if no **binary string in the code is a prefix of any other binary string** in the code. Example:
  - ‘a’=000, ‘b’=0010, ‘c’=0011, ‘d’=01, ‘e’=10, ‘f’ =11
  - 01000100010 = “daeb”
  - Called a **binary prefix code**.



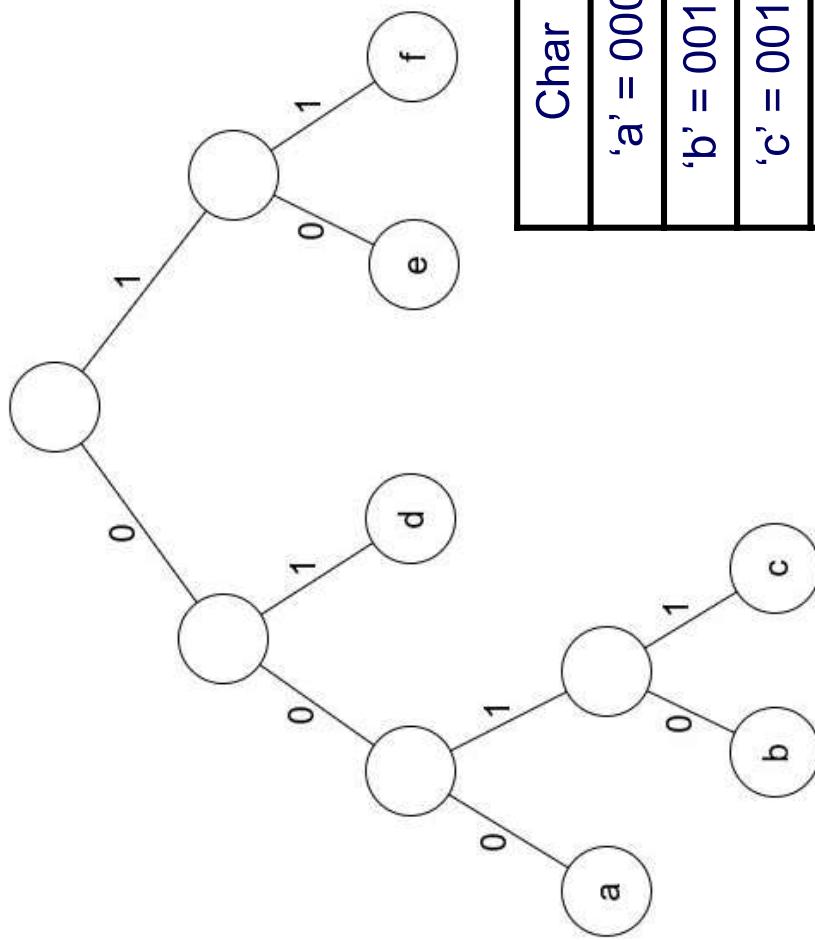
## Representing Binary Prefix Codes

- Binary prefix codes are represented using 2-trees (a binary tree where every non-leaf node has exactly two children).
  - At every internal node, the left branch is labelled 0, and the right branch 1.
- Leaf nodes represent characters.
  - The prefix code of a character is the edge-label sequence along the path from the root to the leaf node containing the character.

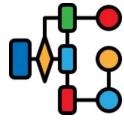


## Example of a Binary Prefix Code

Characters are at the leaves of the tree.

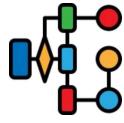


Char	Frequency	No. of bits
'a' = 000	9	27
'b' = 0010	8	32
'c' = 0011	15	60
'd' = 01	3	6
'e' = 10	5	10
'f' = 11	2	4
<b>TOTAL</b>	<b>42 characters</b>	<b>139 bits</b>



## The Encoding Problem

- Given a frequency table for characters in a document, find the/an **optimum** binary-prefix coding tree.
  - Optimum means *minimum length encoded document*.
- Clearly we require short encodings for frequent characters, longer encodings for infrequent ones.
- Example on previous slide is clearly *not* optimal.
- It's worse than a fixed-length 3-bit coding
  - 3 bits is sufficient since we only have 6 characters
    - $3 * 42 = 126$  bits



## Greedy, Take #1

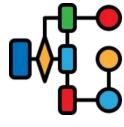
Let's greedily give the most frequent character the shortest possible code (a single 0) then allocate the next available code (10) to the next most frequent character, and so on.

Gives:

Char	Frequency	No. of bits
'c' = 0	15	15
'a' = 10	9	18
'b' = 110	8	24
'e' = 1110	5	20
'd' = 11110	3	15
'f' = 11111	2	10
<b>TOTAL</b>	<b>42 characters</b>	<b>102 bits</b>

Exercise:  
Draw the tree

Better, but is it optimal? Can you find a failing case?



## Greedy, Take #2: Huffman Coding

- Let's greedily give the least common characters the longest encodings!

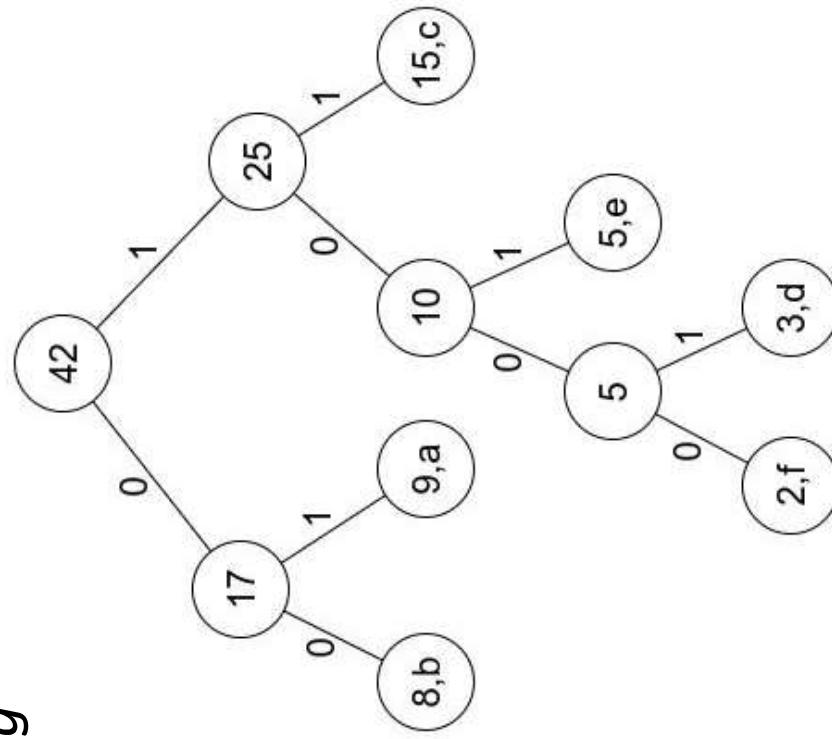
- An algorithm called *Huffman coding*

- It generates ***optimal*** binary prefix codes for a given set of frequencies.

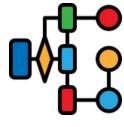
- Solution for previous example:

- Leaf nodes are labelled with (frequency, char)

- Internal nodes are labelled with total frequency of all leaves in their subtree



31



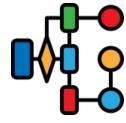
## Huffman Coding algorithm

- Uses a “bottom-up” greedy approach.
- Maintains a list of partial coding trees, sorted in order of total frequency of all characters in the tree
  - Actually we don’t need a fully sorted list; a min-heap will do
  - Start with each character as a trivial tree (just a leaf)
- Repeatedly combine the two trees of lowest frequency

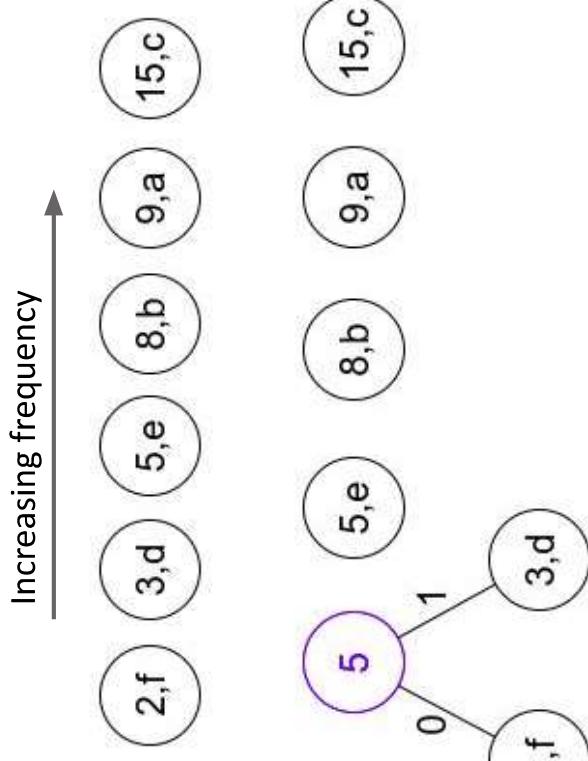
```

trees = MinHeap()
for each character in document:
  trees.insert(Leaf(frequency, character))
while length(trees) > 1:
  left = trees.pop_min()
  right = trees.pop_min()
  trees.insert(Node(left.freq + right.freq, left, right))
    
```

# Huffman Coding - Example



Char	Freq
'a'	9
'b'	8
'c'	15
'd'	3
'e'	5
'f'	2

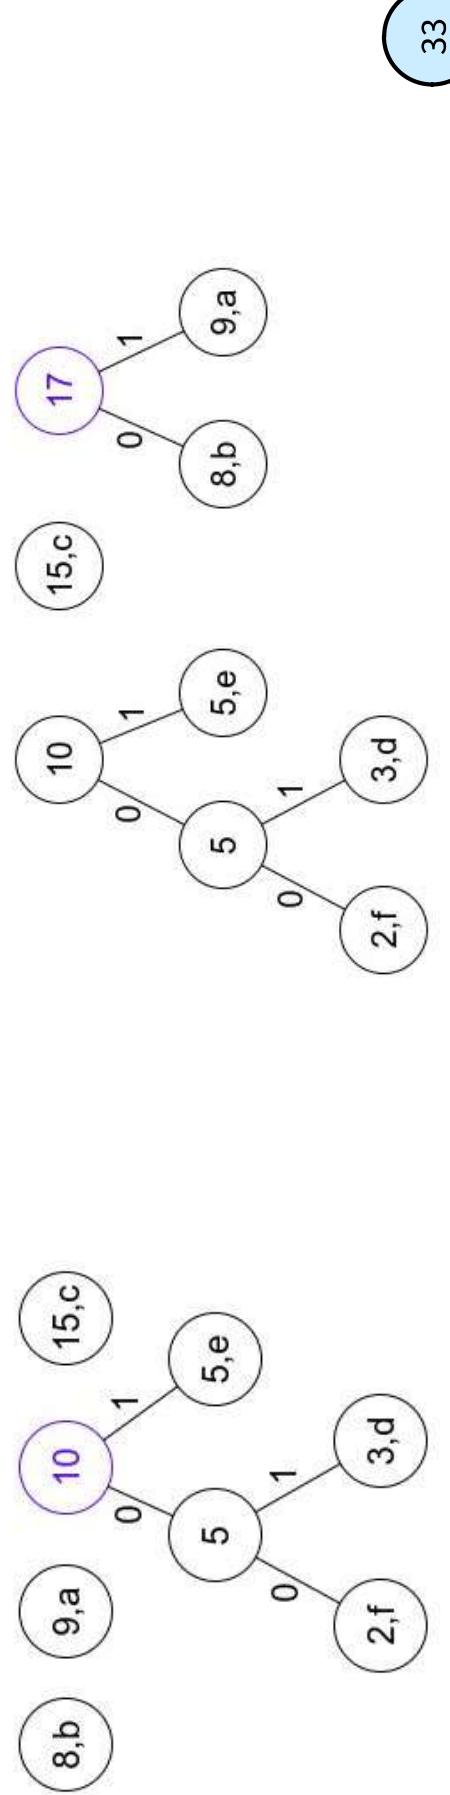


**Stage 1:**

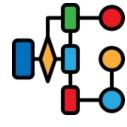
**Stage 2:**

**Stage 3:**

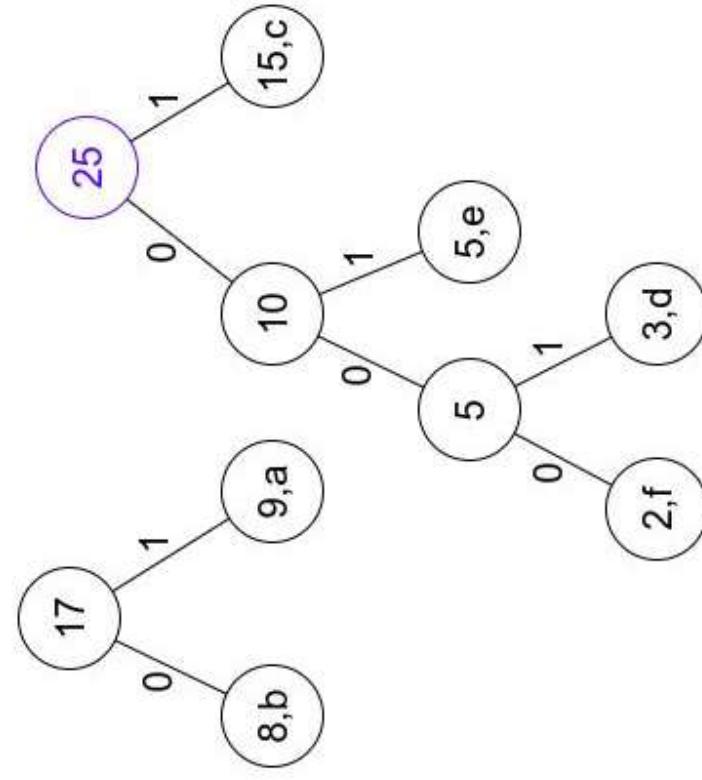
**Stage 4:**



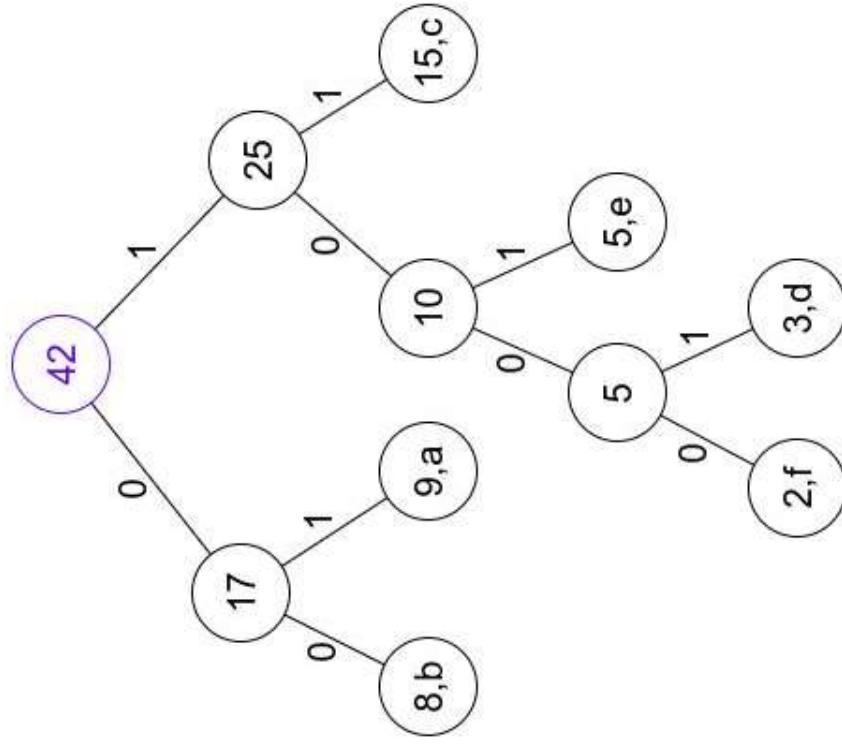
# Huffman Code - Example (cont'd)

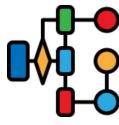


**Stage 5:**



**Stage 6:**

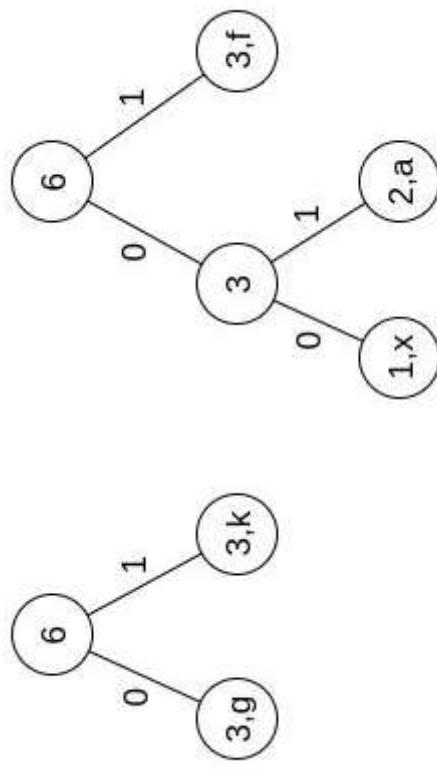


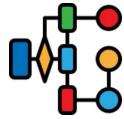


## Ordering trees when frequencies are equal

To ensure a unique solution in lab exercises we require that if multiple trees have the same frequency, they are ordered by the alphabetically smallest leaf node character they contain.

For example: in the image below, both trees have a frequency of 6 but the tree on the right, which has the character 'a' as its smallest leaf, should be sorted to come *before* the tree on the left, which has 'g' as its smallest leaf.

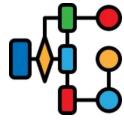




# Huffman Code - Example

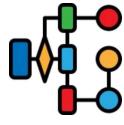
Char	Frequency	No. of bits
'a' = 01	9	18
'b' = 00	8	16
'c' = 11	15	30
'd' = 1001	3	12
'e' = 101	5	15
'f' = 1000	2	8
<b>TOTAL</b>	<b>42 characters</b>	<b>99 bits</b>

- Requires  $99/42 \approx 2.36$  bits/character
- This is an optimal result (as always, with Huffman coding)
  - For proof see [www.cs.utoronto.ca/~brudno/csc373w09/huffman.pdf](http://www.cs.utoronto.ca/~brudno/csc373w09/huffman.pdf)



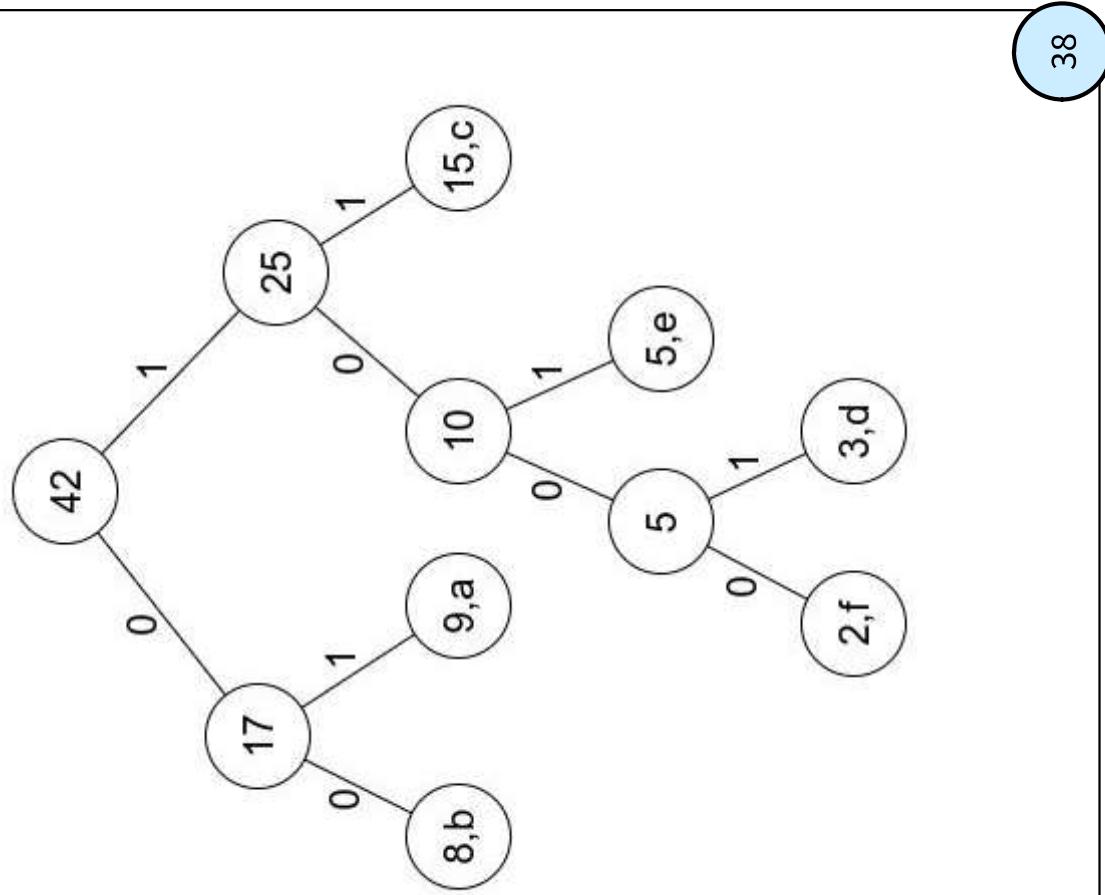
## Huffman Coding: complexity

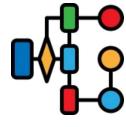
- The *trees* list must support the priority queue operations ‘insert’, and ‘pop\_min’
  - Each is  $O(\log n)$  using a min-heap.
- At each iteration, two elements are removed from the priority queue, and the combined tree’s root is added back to the list. The size reduces by one, and hence the total number of iterations is  $O(n)$ .
- Therefore the total time complexity is  $O(n \log n)$ .



## Decoding Huffman Codes

- To decode a stream of bits, start at the root of the encoding tree, and follow a left-branch for a 0, a right branch for a 1.
- When you reach a leaf, write the character stored at the leaf, and start again at the top of the tree.
- E.g. 0110011100  
= adcb

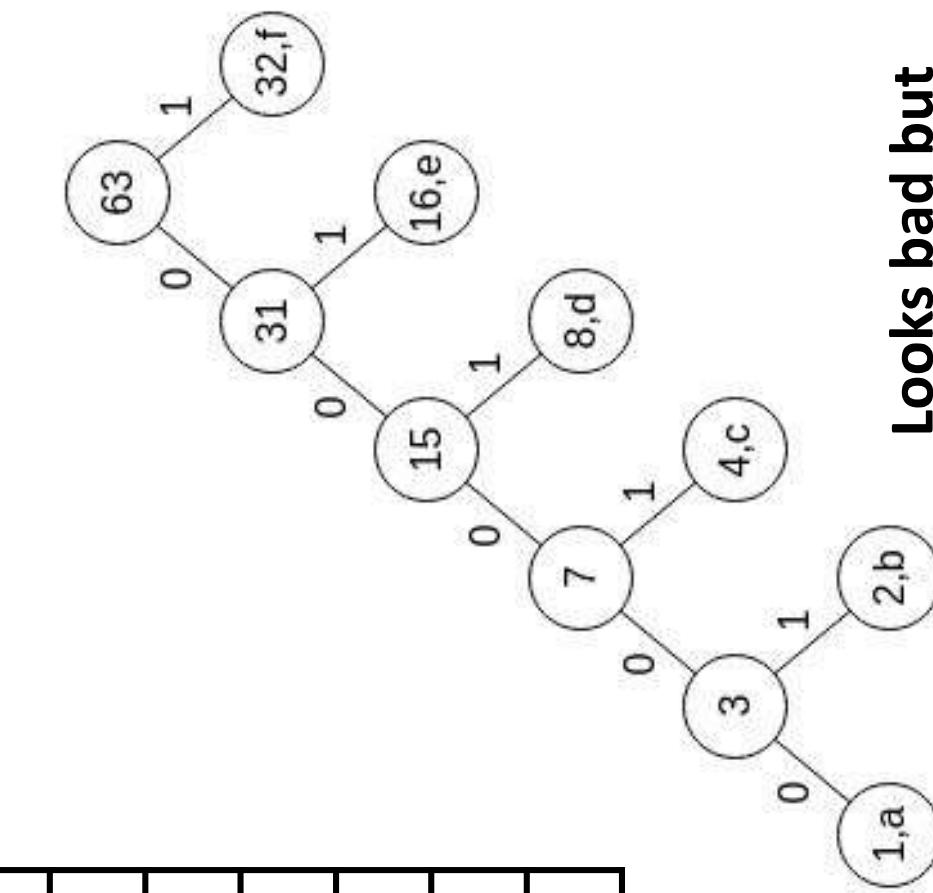




## A Skewed Encoding Tree

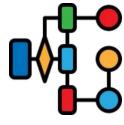
Certain frequency tables can generate unbalanced binary trees

Char	Freq
'a'	1
'b'	2
'c'	4
'd'	8
'e'	16
'f'	32



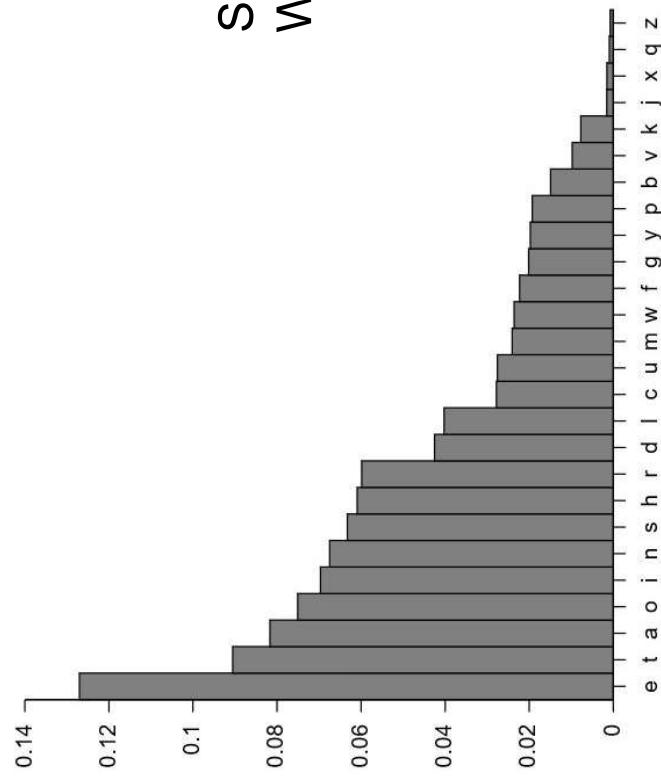
Char	Freq	Enc Bits	Total Bits
'a'	1	5	5
'b'	2	5	10
'c'	4	4	16
'd'	8	3	24
'e'	16	2	32
'f'	32	1	32
Total no. of bits		119	
Avg. no. of bits per char		1.8	

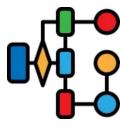
Looks bad but  
still optimal.



## Caveat

- The encoding information must be transmitted along with the text to the decoding system.
  - An added cost
- OR: the frequency information for generic English text (say) may be used, e.g.





## Advantages/disadvantages

Huffman coding gives optimal encoding based on *global* character frequencies only, but:

- Whole document must be processed first.
- Can't do local optimisation (e.g. change in char frequencies half way through a long document).
- Encoding table must itself be encoded into the document.
- Can't recognise *patterns*, e.g. common words.

Lempel-Ziv algorithm (and derivatives like LZW) generally give better compression by using a dynamic compression table with common substring recognition.

*Deflate* algorithm (<https://en.wikipedia.org/wiki/DEFLATE>) combines LZ and Huffman coding.

- Used by zip, pkzip, gzip, png, etc